

# Contents

Self-RAG Implementation Report .....	2
1. Overview .....	2
2. Objectives .....	2
3. High-Level Architecture.....	3
Core Components .....	3
4. Graph State Design .....	3
State Fields .....	3
5. Node-Level Description.....	3
5.1 Retrieve Node.....	3
5.2 Document Relevance Grading Node .....	4
5.3 Generate Node .....	4
5.4 Query Transformation Node .....	4
6. Control Flow Logic .....	4
7. Alignment with Self-RAG Framework .....	4
Architectural Differences.....	5
9. Limitations .....	5
10. Conclusion.....	5

# Self-RAG Implementation Report

## 1. Overview

This report documents the design, logic, and architecture of the Self-RAG (Self-Reflective Retrieval-Augmented Generation) system implemented using LangGraph. The system is designed to improve answer quality by iteratively retrieving documents, generating answers, and self-evaluating relevance, groundedness, and usefulness before producing a final response.

The implementation follows the core ideas proposed in the Self-RAG framework while adapting them for a production-oriented environment using HuggingFace models and a local document corpus.

## 2. Objectives

The primary goals of this implementation are:

- Reduce hallucinations in RAG-based systems
- Ensure retrieved documents are relevant to the query
- Verify that generated answers are grounded in retrieved content
- Ensure answers directly address the user's question
- Automatically recover from failure cases using query rewriting and regeneration

### 3. High-Level Architecture

The system is implemented as a stateful directed graph using LangGraph. Each node performs a specific operation, and conditional edges determine the control flow based on intermediate evaluation results.

#### Core Components

- **Retriever:** Fetches top-k relevant document chunks from a Chroma vector store
- **LLM Generator:** Produces answers using retrieved context
- **Self-Grading Modules:**
  - Document relevance grader
  - Groundedness (hallucination) checker
  - Answer usefulness checker
- **Query Rewriter:** Reformulates queries when retrieval or generation fails

### 4. Graph State Design

The system maintains a shared mutable state across nodes:

```
class GraphState(TypedDict):
    question: str
    documents: list
    generation: str
    steps: List[str]
```

#### State Fields

- **question:** Current user question (may be rewritten)
- **documents:** Retrieved and filtered document chunks
- **generation:** Latest LLM-generated answer
- **steps:** Execution trace for debugging and observability

### 5. Node-Level Description

#### 5.1 Retrieve Node

- Retrieves candidate documents using vector similarity search
- Adds retrieval logs to the execution trace

**Purpose:** Provide candidate evidence for answer generation

## 5.2 Document Relevance Grading Node

- Uses an LLM-based binary classifier
- Filters out documents not relevant to the question

**Decision Outcome:** - If no relevant documents remain → trigger query rewriting - Else → proceed to answer generation

## 5.3 Generate Node

- Generates an answer using retrieved documents
- Immediately evaluates groundedness and usefulness

**Possible Outcomes:** - Hallucinated answer → regenerate - Grounded but not useful → rewrite query - Grounded and useful → terminate successfully

## 5.4 Query Transformation Node

- Rewrites the original question to improve retrieval quality
- Clears document state before re-retrieval

**Purpose:** Recovery from poor retrieval or misaligned answers

# 6. Control Flow Logic

The graph follows this execution loop:

1. Start → Retrieve
2. Retrieve → Grade Documents
3. If no relevant docs → Rewrite Query → Retrieve
4. If relevant docs → Generate
5. If hallucinated → Regenerate
6. If not useful → Rewrite Query → Retrieve
7. If grounded & useful → End

This loop ensures robustness and self-correction.

# 7. Alignment with Self-RAG Framework

The implementation follows the Self-RAG design:

- Retrieval decision-making
- Document relevance validation
- Groundedness verification
- Answer usefulness scoring
- Iterative refinement loop

## Architectural Differences

- Uses HuggingFace (Mistral-7B-Instruct) instead of OpenAI models
- Consolidates some grading logic within the generation node
- Operates on a local enterprise document corpus

## 9. Limitations

- LLM-based graders may occasionally misclassify
- Increased latency due to multiple LLM calls (40s+ for a unit test)
- Requires careful prompt tuning for graders

## 10. Conclusion

This Self-RAG LangGraph implementation demonstrates a robust, self-correcting RAG pipeline suitable for enterprise knowledge systems. While architecturally adapted for practical constraints, it faithfully implements the core Self-RAG reasoning paradigm and significantly improves answer reliability and relevance.